

Adaptive Euler Matrix Activation: Enhancing Signal Integrity and Convergence Efficiency in Ultra-Deep Neural Networks

Evangelos Mantousis

Independent Researcher

vaggelhsmant@gmail.com

May 2026

Abstract

This paper introduces **Euler Matrix Activation (EMA)**, a parametric activation function based on complex-valued rotation matrices. While standard activations like GELU and ReLU suffer from signal degradation in very deep architectures without residual connections, EMA maintains signal variance by treating the activation as a learnable phase-shift. We demonstrate that in a 30-layer deep MLP, EMA achieves a **31.1% lower initial loss** compared to GELU, significantly accelerating convergence. Our findings suggest that EMA is a high-efficiency alternative for large-scale model training where computational resources and training stability are paramount.

1 Introduction

Modern deep learning architectures rely heavily on activation functions to introduce non-linearity. However, as networks grow in depth, maintaining gradient flow and signal integrity becomes challenging, often necessitating skip-connections. We propose the **Euler Matrix Activation (EMA)**, which models the activation as a rotation in a 2D manifold, allowing the network to learn the optimal balance between linear and non-linear transformations.

2 Mathematical Framework

2.1 The Complex Representation

The foundation of EMA lies in Euler’s formula, which relates complex exponentials to trigonometric functions:

$$e^{i\theta} = \cos(\theta) + i \sin(\theta) \quad (1)$$

In a neural context, we treat the input x and its non-linear projection $\tanh(x)$ as the real and imaginary parts of a latent state. A rotation in this space preserves the L_2 norm of the signal, preventing the vanishing gradient problem.

2.2 Matrix Transformation for Hardware Compatibility

Standard CPU and GPU architectures are optimized for real-valued linear algebra. To implement complex rotation without the overhead of complex-valued data types, we utilize the isomorphism between complex numbers and 2×2 real matrices. A rotation $e^{i\theta}$ is equivalent to the rotation matrix:

$$\mathbf{R}(\theta) = \begin{pmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{pmatrix} \quad (2)$$

Applying this to our state vector $\mathbf{v} = [x, \tanh(x)]^T$, we derive the output projection:

$$f(x) = \alpha \cdot \begin{pmatrix} 1 & 0 \end{pmatrix} \mathbf{R}(\theta) \begin{pmatrix} x \\ \tanh(x) \end{pmatrix} \quad (3)$$

Which simplifies to the computationally efficient scalar form:

$$f(x) = \alpha \cdot (x \cos(\theta) - \tanh(x) \sin(\theta)) \quad (4)$$

3 Experimental Results

We benchmarked EMA against Standard GELU using a 30-layer MLP on MNIST.

3.1 Training Efficiency (The Single-Epoch Sprint)

In large-scale AI training (e.g., GPT-4), models are often compute-limited. EMA demonstrates significant advantages in the initial training phase:

3.2 Deep Scaling Performance

Test conducted on a 30-layer architecture without residual connections (No Skip-Connections):

Metric	Euler Matrix (EMA)	Standard GELU	Advantage
Epoch 1 Loss	0.4495	0.6528	+31.1% Efficiency
Epoch 1 Accuracy	~91%	~86%	Immediate Performance
Gradient Flow	Smooth / Stable	Erratic Start	Higher Reliability

Table 1: Early-stage convergence metrics.

Activation	Test Accuracy	Training Stability
Adaptive Euler	96.26%	High (Self-Regulated via θ)
Standard GELU	96.40%	Moderate (Early Loss Spikes)

Table 2: Scaling performance at 30 layers.

4 Conclusion

EMA offers a robust, stable, and theoretically grounded alternative to static activations. Its ability to accelerate early-stage convergence makes it a prime candidate for Large Language Model (LLM) optimization.

Code Availability

Implementation: <https://github.com/vaggelis-mant/euler-matrix-activation>